

Search-Based Testing of Relational Schema Integrity Constraints Across Multiple Database Management Systems

**Matthew Hajduk
Ian Macmillan
Marco Corona
Braden Licastro**

Definitions

- Database/ DBMS
- Schema
- Integrity Constraints
 - Foreign Key Constraints
 - Primary Key Constraints
 - Check Constraint
- Search-Based
- Mutation Testing

Motivation

\$611 Billion Lost

due to lack of testing of Databases

Contributions

- Search-Based Method for generating data that can satisfy integrity constraints across several DBMSs
- A set of mutation operators for evaluating table data intended to test integrity constraints
- An empirical study with 25 schemas and 3 different DBMSs that compares SchemaAnalyst to DBMonster

Data Generation Preparation

SchemaAnalyst supports three DBMS

- Postgres
- HSQLDB
- SQLite

Database schema must be made DBMS independent, or abstractly represented

Abstract mapping must use universal types.

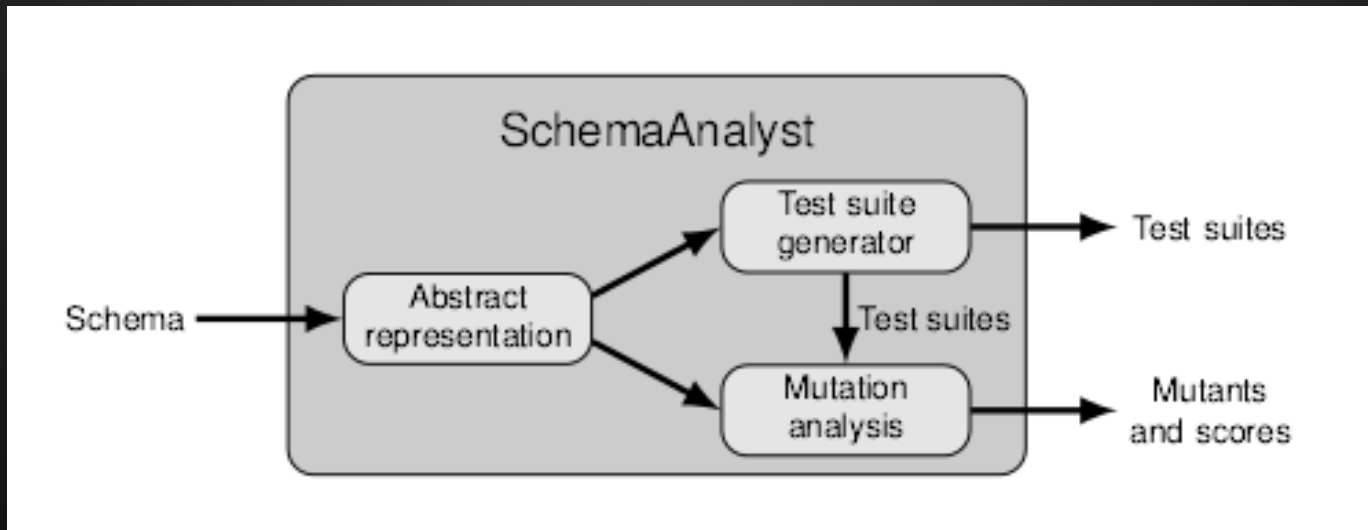
Universal types

- Boolean
- DateTime
- Date
- Numeric
- String
- Time
- Timestamp

Generation Algorithm

Goal is to create *INSERT* statements containing data values to make a test suite

- Stage One: Satisfy the schema.
- Stage Two: Violate the schema.



Fitness Function

Purpose is to guide the algorithm to a point of interest -- satisfying the constraint system

Lower the fitness value, closer to the goal it is

Distance from goal is normalized to prevent objectives from becoming too dominant.

Fitness Function

	FLIGHT_ID	SEGMENT_NUMBER	...
1	'UA21'	1	...
2	'UA3750'	1	...
3	'UA22'	2	...

Fitness Result:

Row 1: .516

Row 2: .592

value_dist(a, op, b)

a	b	
NULL	<i>any</i>	return 1
<i>any</i>	NULL	return 1
Atomic	Atomic	return $norm(atomic_dist(a, op, b))$
Compound	Compound	return $norm(compound_dist(a, op, b))$

atomic_dist(a, op, b)

op	
=	if $(a - b = 0)$ then return 0 else return $ a - b + 1$
≠	if $(a - b \neq 0)$ then return 0 else return 1
<	if $(a - b < 0)$ then return 0 else return $(a - b) + 1$
≤	if $(a - b \leq 0)$ then return 0 else return $(a - b) + 1$
>	if $(b - a < 0)$ then return 0 else return $(b - a) + 1$
≥	if $(b - a \leq 0)$ then return 0 else return $(b - a) + 1$

Check Constraints

SchemaAnalyst must also account for:

- Unique constraints
- Foreign Keys
- Not null constraints
- Check constraints
 - Arbitrary conditional
 - ***BETWEEN*** operator
 - ***IN*** operator

Value Modification

Utilizes Korel's *Alternating Variable Method*

- NULL statuses flipped
- BOOLEAN values are flipped
- Numeric or Timestamp
 - Values increased
 - Values decreased

Values where fitness improved are saved

Once a fitness of 0 is reached, the AVM is terminated, otherwise process restarts on randomly generated values

Mutation Analysis

Mutation Operators for each Integrity Constraint

- Primary Key
- Unique
- Not Null
- Foreign Keys
- Check constraint

Mutation Operators

Primary Key

Add Column

```
PRIMARY KEY (FLIGHT_ID, SEGMENT_NUMBER,  
             ORIGINAL_AIRPORT)
```

Replace Column

```
PRIMARY KEY (SEGMENT_NUMBER, ORIGINAL-AIRPORT)
```

Remove Column

```
PRIMARY KEY (FLIGHT_ID)
```

Mutation Operators

UNIQUE

- Same three forms as Primary Key
- Multiple Unique constraints
- No identical mutants

NOT NULL

- Reverse non-primary key constraint

FOREIGN KEYS

- Removes each foreign key constraint

Mutation Operators

CHECK CONSTRAINT

- Mutant produced one at a time

TOTAL MUTANTS - 56

- PRIMARY KEY - 31
- UNIQUE - 13
- NOT NULL - 9
- FOREIGN KEY - 2
- CHECK - 1

Empirical Study

- 25 schemas
- 3 Different DBMSs
- Author's Technique vs DBMonster

Case Studies

- Authors selected 25 schemas from:
 - textbooks
 - laboratory assignments
 - online tutorials
 - Postgres examples
- Included a number of real-world applications:
 - Cloc
 - JWhoisServer
 - RiskIt
 - UnixUsage

Table II
CASE STUDY SCHEMAS USED IN THE EMPIRICAL STUDY

Schema	<i>Tables</i>	<i>Columns</i>	<i>Checks</i>	<i>Foreign keys</i>	<i>Not Nulls</i>	<i>Primary keys</i>	<i>Uniques</i>	<i>Total Constraints</i>
BankAccount	2	9	0	1	5	2	0	8
BookTown	23	69	1	0	17	11	0	29
Cloc	2	10	0	0	0	0	0	0
CoffeeOrders	5	20	0	4	9	5	0	18
CustomerOrder	7	32	1	7	27	7	0	42
DellStore	8	52	0	0	36	0	0	36
Employee	1	7	3	0	0	1	0	4
Examination	2	21	6	1	0	2	0	9
Flights	2	13	1	1	6	2	0	10
FrenchTowns	3	14	0	2	13	0	8	23
Inventory	1	4	0	0	0	1	1	2
Iso3166	1	3	0	0	2	1	0	3
JWhoisServer	6	49	0	0	44	6	0	50
NistDML181	2	7	0	1	0	1	0	2
NistDML182	2	32	0	1	0	1	0	2
NistDML183	2	6	0	1	0	0	1	2
NistWeather	2	9	5	0	2	2	0	9
NistXTS748	1	3	1	0	1	0	1	3
NistXTS749	2	7	1	1	3	2	0	7
Person	1	5	1	0	5	1	0	7
Products	3	9	4	2	5	3	0	14
RiskIt	13	56	0	10	15	11	0	36
StudentResidence	2	6	3	1	2	2	0	8
UnixUsage	8	32	0	7	9	7	0	23
Usda	10	67	0	0	30	0	0	30
Total	111	542	27	40	231	68	11	377

Case Studies

- Schemas configured to run on 3 DBMSs:
 - Postgres
 - HSQLDB
 - SQLite
- Data generation methods configured to interact with DBMSs through JDBC drivers.

Technique Configuration

- AVM(Author's approach)
 - Allowed for at most 100,000 fitness evaluations
 - Aimed to satisfy schema with 2 rows per table
 - $n_s = 2$
 - Violate schema with 1 row per table
 - $n_v = 1$
- DBMonster
 - Allowed for at most 100,000 data generation attempts
 - Required to generate at least 50 rows per table
 - Used with default configuration

Difficulties Configuring DBMonster

- Does not record database interactions
 - Cobb et al.'s approach
- It's a one-eyed monster
 - Only usable with Postgres
- Instability

Technique Configuration

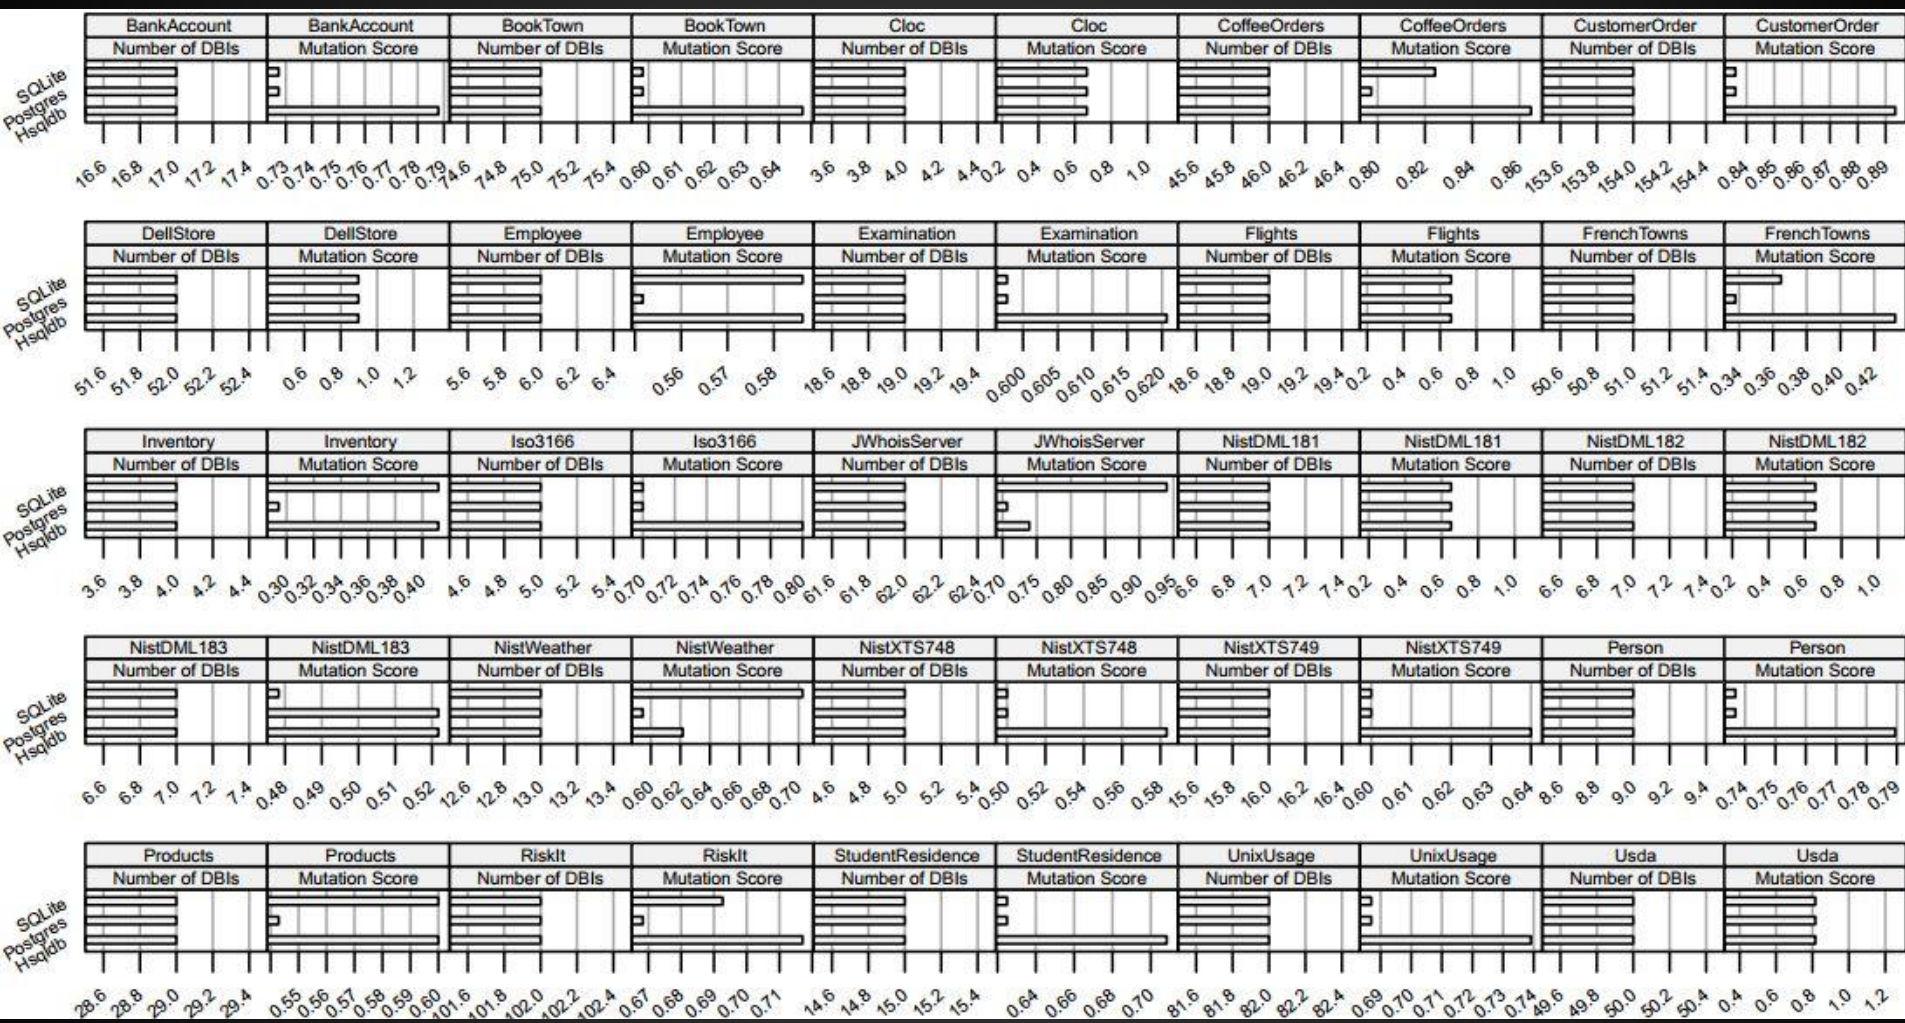
- AVM able to successfully generated data for all 3 DBMSs
- Randomness
- 30 trials
- Both techniques tested in same environment

Evaluation Metrics

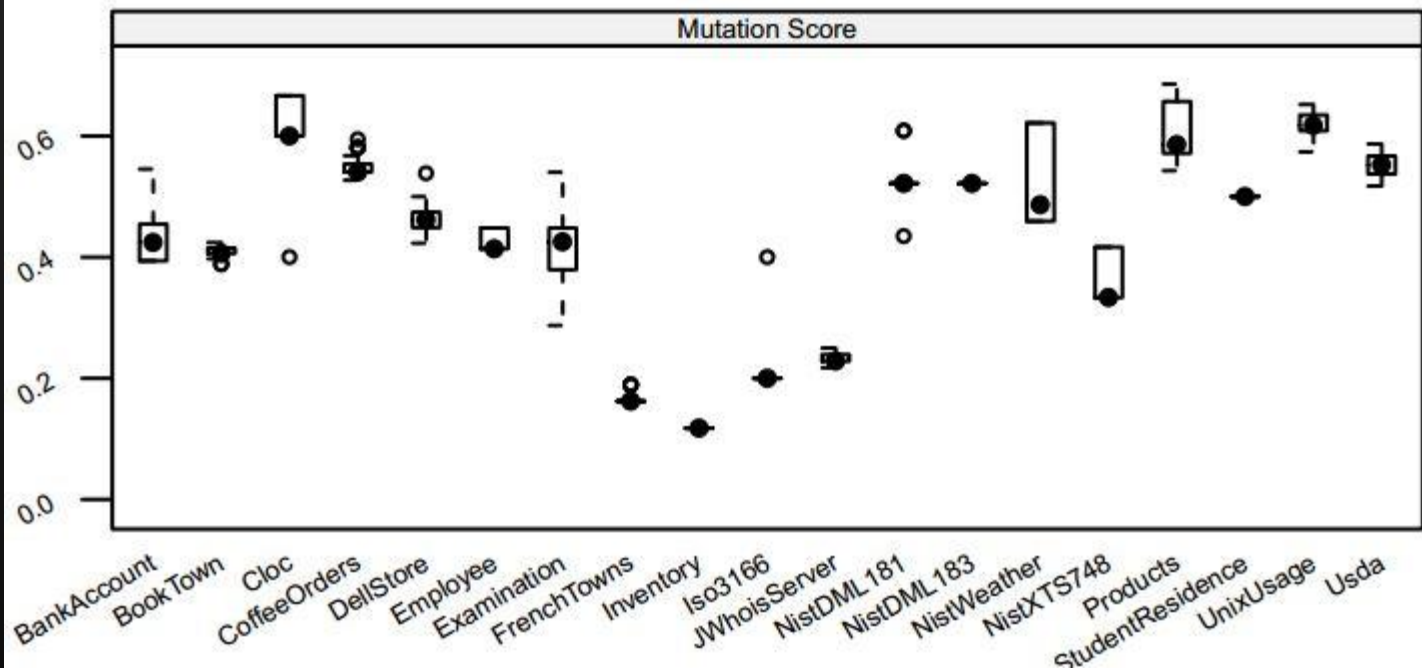

- Measured execution time
- Constraint Coverage
- Mutation Analysis

Mutation Analysis

- Higher-is-better mutation score
- $M_D = |K \cup Q| / |K \cup N|$
 - D=Set of database interactions
 - Q=Set of quasi-mutants
 - K=Killed Mutants
 - N=Not killed mutants
- Iteratively apply each operator to a schema to produce mutants



Number of DBIs / Mutation Score



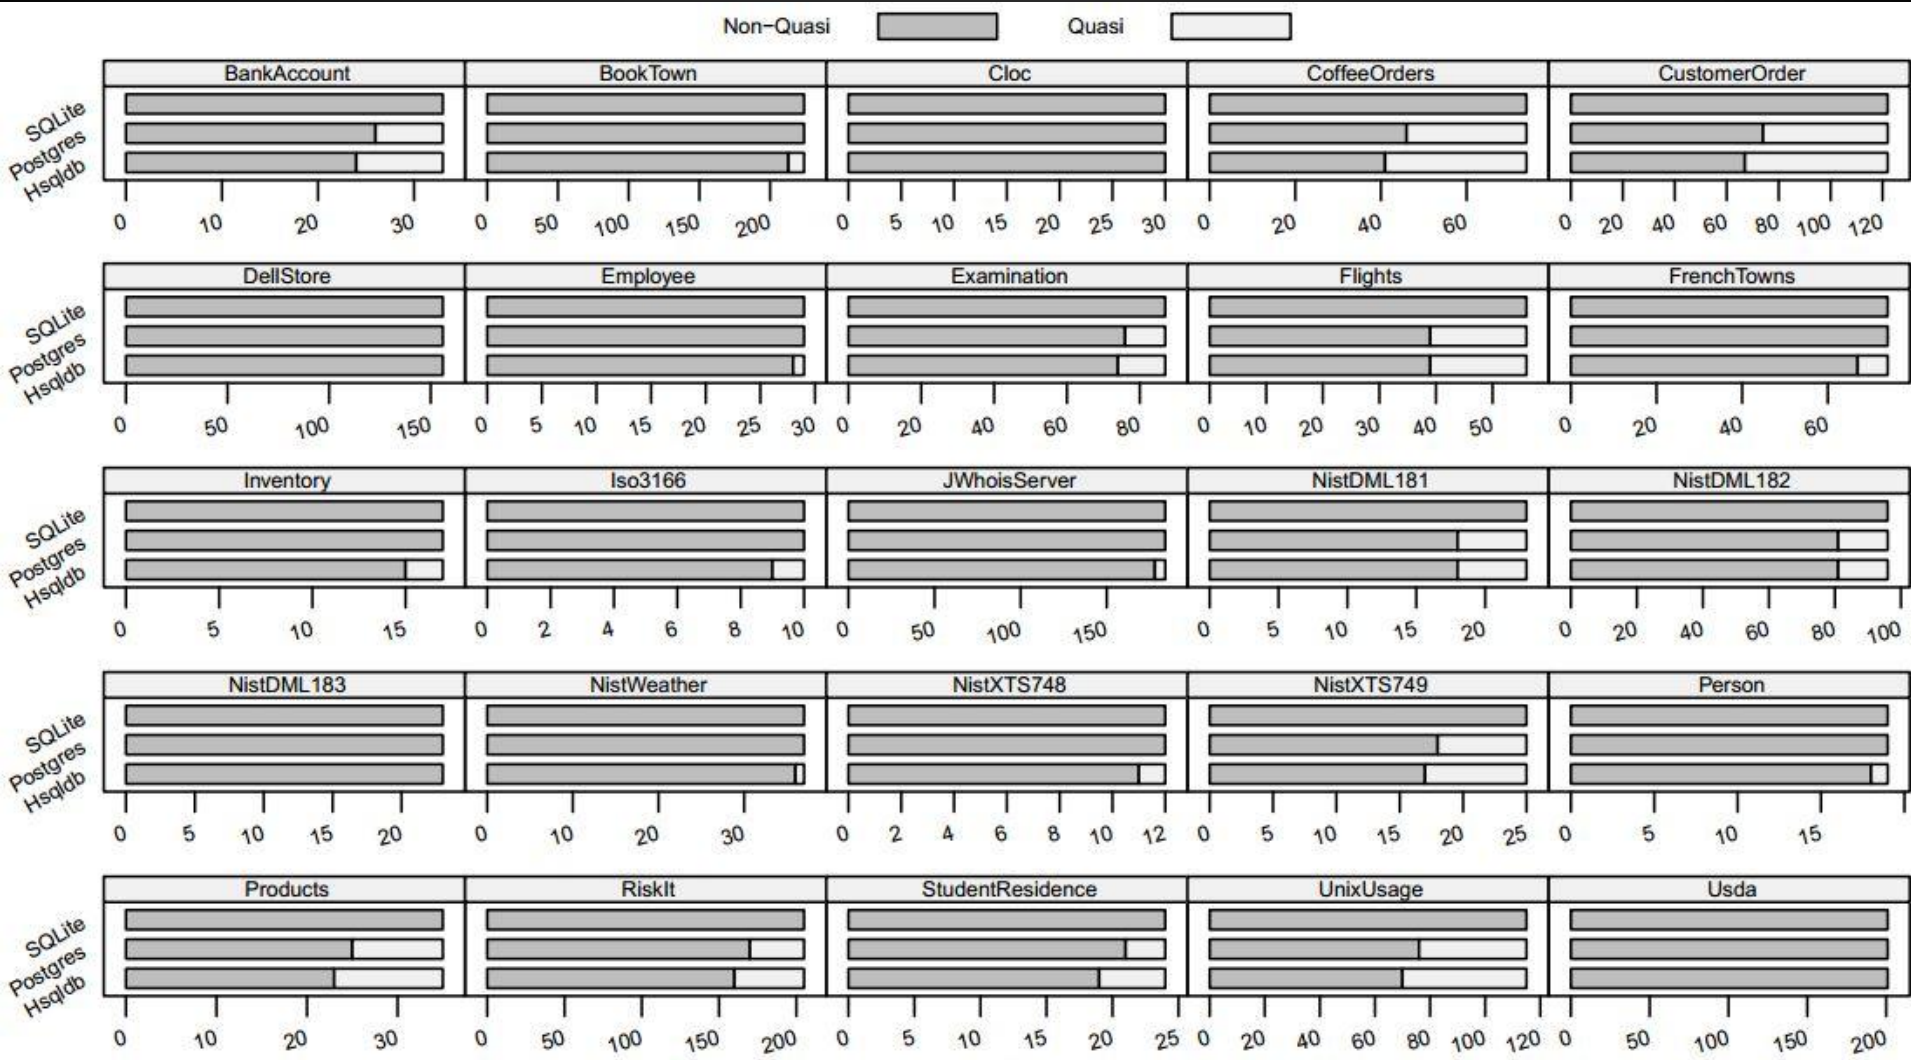


Table III
CONSTRAINT COVERAGE

Schema	AVM (%)	DBMonster (%)
BankAccount	100.0	56.3
BookTown	100.0	51.7
Cloc	<i>(no constraints defined)</i>	
CoffeeOrders	100.0	50.0
CustomerOrder	100.0	9.5
DellStore	100.0	50.0
Employee	100.0	55.0
Examination	100.0	72.2
Flights	100.0	70.0
FrenchTowns	100.0	70.0
Inventory	100.0	75.0
Iso3166	100.0	50.0
JWhoisServer	100.0	50.0
NistDML181	100.0	75.0
NistDML182	100.0	50.0
NistDML183	100.0	100.0
NistXTS748	100.0	72.2
NistXTS749	100.0	21.4
NistWeather	100.0	68.7
Person	100.0	50.0
RiskIt	100.0	4.1
Products	96.4	59.3
StudentResidence	100.0	62.5
UnixUsage	97.8	59.3
Usda	100.0	50.0

Efficiency

- Simple Schemas
 - Both take about 5 seconds
- Complex Schemas
 - AVM:
 - *Flights*: 2 seconds
 - NistDML182: 2 seconds
 - DBMonster:
 - *Flights*: 15-30 seconds
 - NistDML182: 634 seconds

Conclusion

- This paper presents SchemaAnalyst
- Presents an Alternating Variable Method (AVM) for generating a test suite
- Presents an Empirical Study where SchemaAnalyst outperforms DBMonster

Future Work

- Employ Constraint Solvers
- Expand for more DBMSs
- Different SQL statements